

# CS267 Homework 2 (Part 3)

## Parallelizing a Particle Simulation

Team members: Tzu-Chuan Lin, Chin-An Chen, Byron Hsu

### Abstract

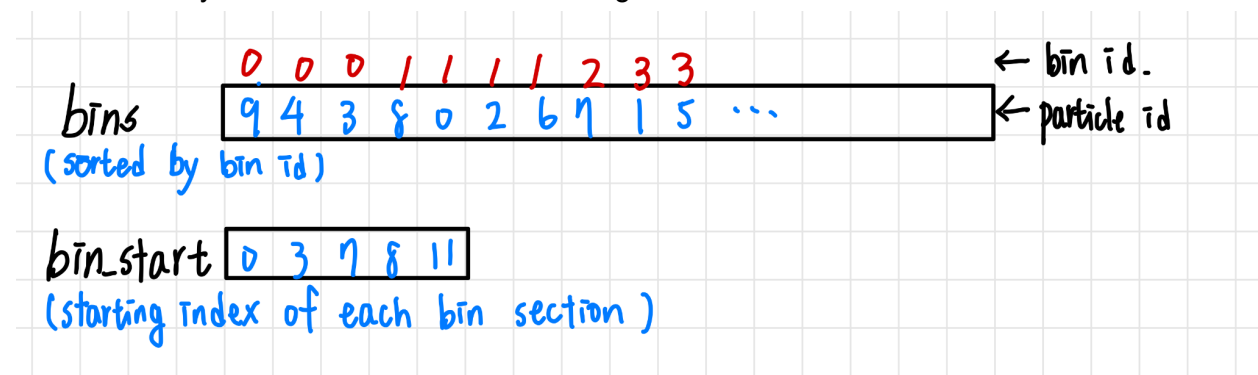
In this assignment, we will be parallelizing a toy particle simulation (similar simulations are used in mechanics, biology, and astronomy) with GPU CUDA. In our simulation, particles interact by repelling one another but only when closer than a cutoff distance.

Suppose we have a code that runs in time  $T = O(n)$  on a single processor. Then we'd hope to run close to time  $T/p$  when using  $p$  processors. You will attempt to reach this speed up with a GPU. As the result, we got the best performance on GPU in around **1.6 seconds for 1M particles**.

### Implementation

As the methods recommended by the GSI, instead of using `std::vector`, which would cost lots of copies from host to GPU per step, we used arrays.

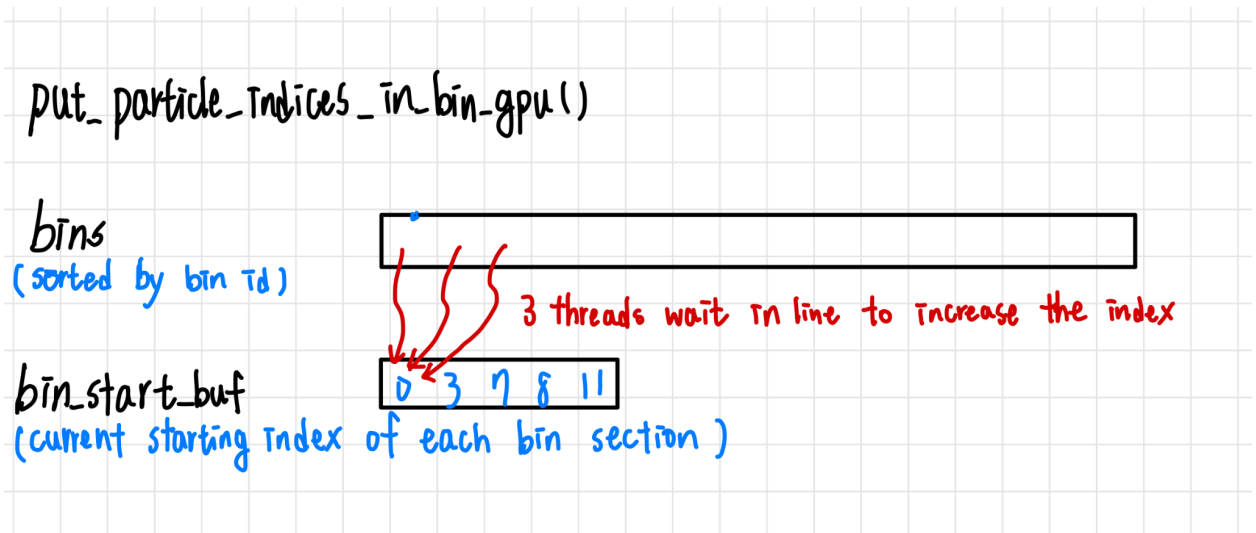
There are arrays as below we used in this assignment.



- `particle_t* parts_gpu`: the array of given particles (size == #of particles)
- `int* bins`: array of particles sorted by bin id (size == #of particles)
- `int* bin_start`: Prefix sum of the bin counts (size == #of bins)
- `int* bin_start_buf`: instead of literally sorting algorithm to update the array `bins`, we used `bin_start_buf` to keep track of the indices on array `bins` we are going to modify.

## Dataflow:

- `init_simulation()`: memory allocation on GPU for arrays `bins`, `bin_start`, and `bin_start_buf`
- `simulation_one_step()`:
  - `count_each_bin_gpu()` and `inclusive_scan()`. Calculate the prefix sum for each bin, and record the result in array `bin_start`.

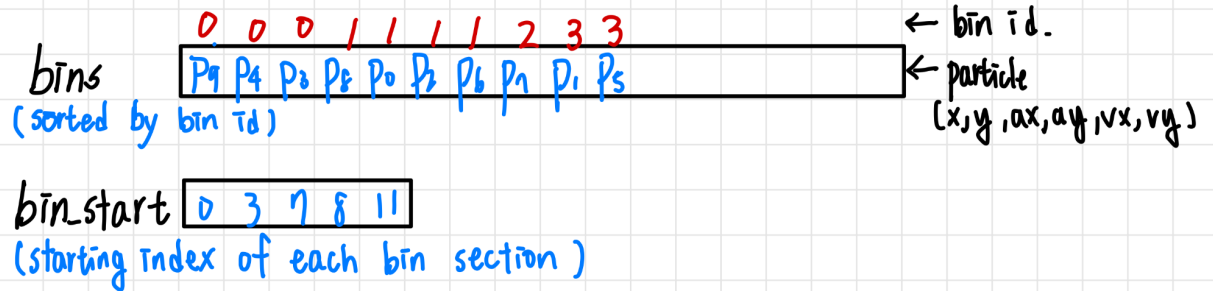


- `put_particle_indices_in_bin_gpu()`: Sort the particles with their bin indices and store them in an array `bin`.
- `clear_ax_ay()`: Initialization of the acceleration of each particle
- `compute_forces_gpu()`: Calculation of the force between particles and their acceleration.
- `move_gpu()`: Move the particles based on the acceleration we got in `compute_forces_gpu`.

## Experiments

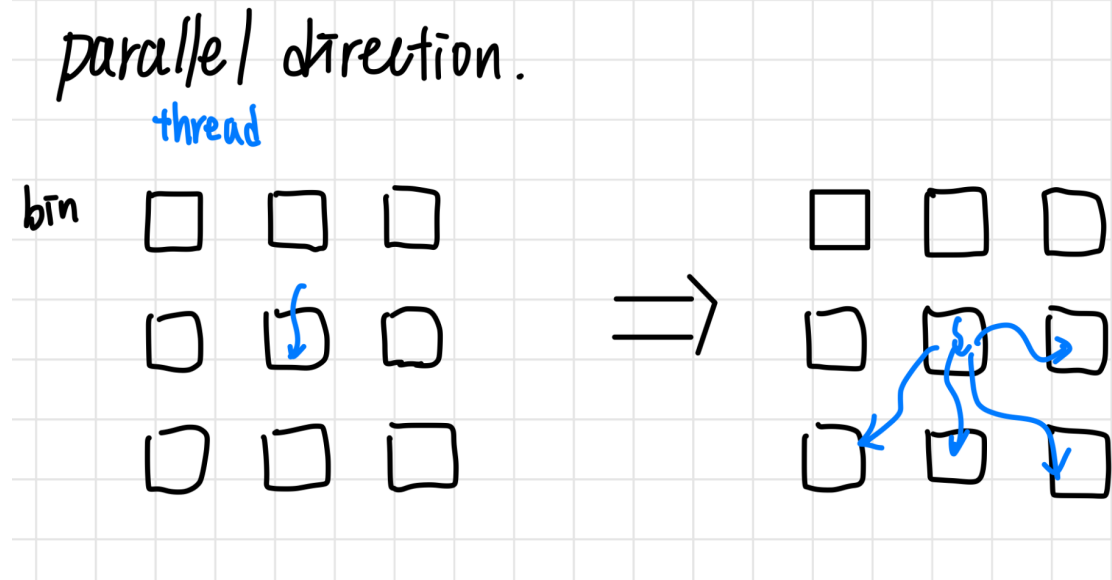
1. One direction force:  
Calculate the force for the current particle only; hence, we **do not need atomic add** when updating the acceleration.
2. Bi-direction force:  
Calculate the forces for the current particle and the particle paired with it; hence, we need **atomic add** when updating the acceleration
3. Bi-direction force calculation with particle copying:  
Not only storing the particle indices in the array `bin`, but we also stored the position, velocity, and acceleration of the particle in the array `bin`.

particle copying.



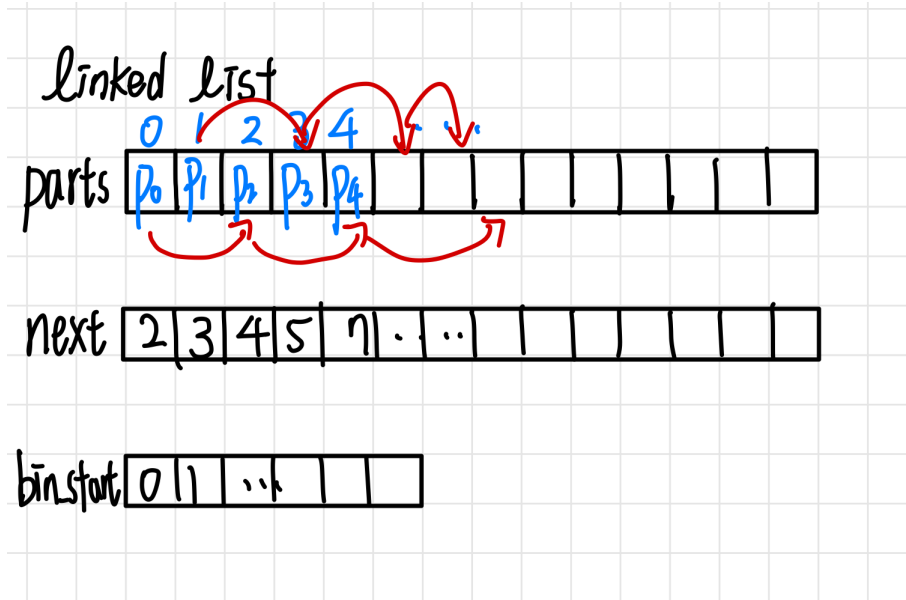
4. Bi-direction force and parallel directions:

Because when we calculate the forces bi-directionally, we have to go through five bins (current, right, bottom left, bottom, bottom right), we then modified the algorithm to calculate the forces **all by once in parallel**.



5. Bi-direction force with a **linked list for each bin**:

We used the array **next** to record the next particle of the corresponding particle in the array **gpu\_parts**. When constructing the next array, we traversed the **gpu\_parts** and logically added the particle at the end of the list (bin).

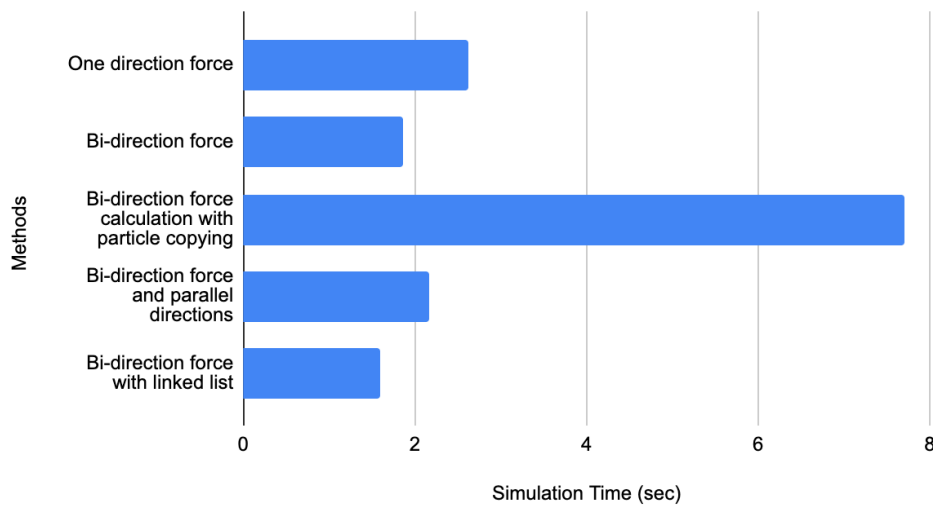


## Design choice and the effects

Tested on 1M particles

Method	1	2	3	4	5
Simulation time (sec)	2.62339	1.86601	7.71509	2.17554	1.59818

Simulation Time (sec) vs. Methods



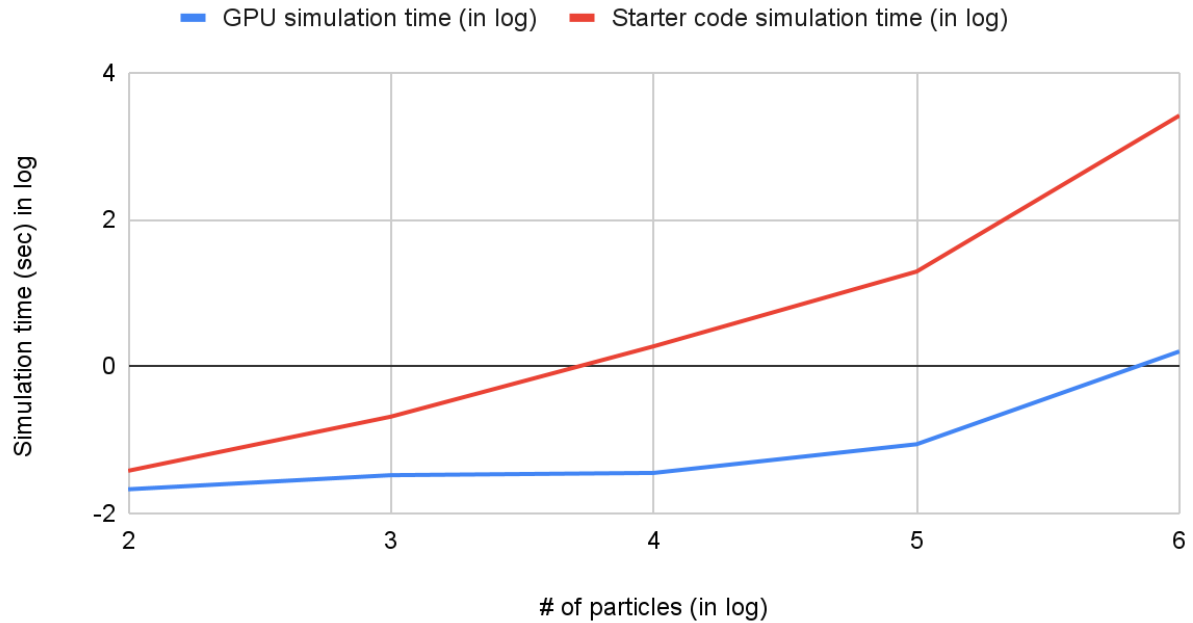
# Performances

Log-log (Time versus particles)

Tested with method 5 (bi-direction and linked list)

# of particles	100	1000	10000	100000	1000000
Starter-code on GPU (O(n^2))	0.0379855	0.207346	1.87929	19.531	2593.46
GPU	0.0212772	0.033097	0.0355539	0.0872949	1.59818

simulation time (in log) v.s. # of particles (in log)



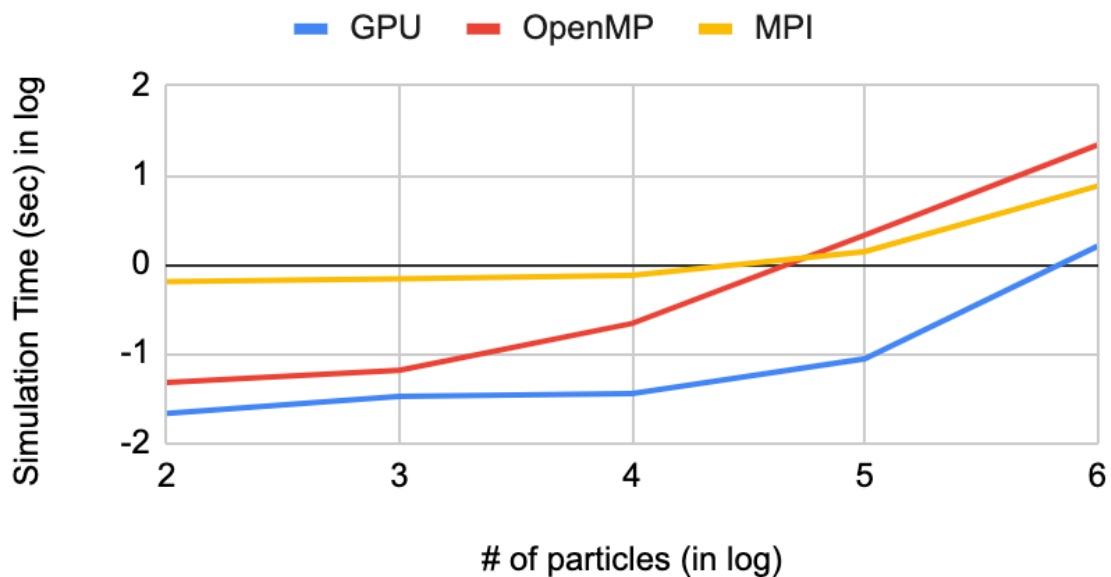
## Parallizing using GPU versus OpenMP and MPI

MPI: 2 node 68 tasks per node

OpenMP: 1 node 68 cores 1 thread per core

	100	1000	10000	100000	1000000
GPU	0.0212772	0.033097	0.0355539	0.0872949	1.59818
OpenMP	0.0475355	0.0648152	0.217363	2.12787	21.8643
MPI	0.640243	0.692009	0.755319	1.39015	7.5557

## GPU, OpenMP and MPI Comparisons



## Breaking down the runtime

- Computation time: slightly in proportion to the size of particles
- Synchronization time: `atomic` / `cudaDeviceSynchronize`
- Communication time: `cudaMemcpy`

## Contribution among the team

Tzu-Chuan: Implementation

Chin-An: Text of the report

Byron: Illustration of the report